

THE HONG KONG
UNIVERSITY OF SCIENCE
AND TECHNOLOGY



ClearAgent: Agentic Binary Analysis for Effective Vulnerability Detection

Xiang Chen, Anshunkang Zhou, Chengfeng Ye, Charles Zhang

The Hong Kong University of Science and Technology

2025-10-15 @ LMPL '25

Contents

1. Background	2
2. ClearAgent	8
3. Preliminary Evaluation	14
4. Future Work	19

I. Background

Binary Code Analysis

I. Background

- Analyzable binary code formats: Assembly → IR → Source code

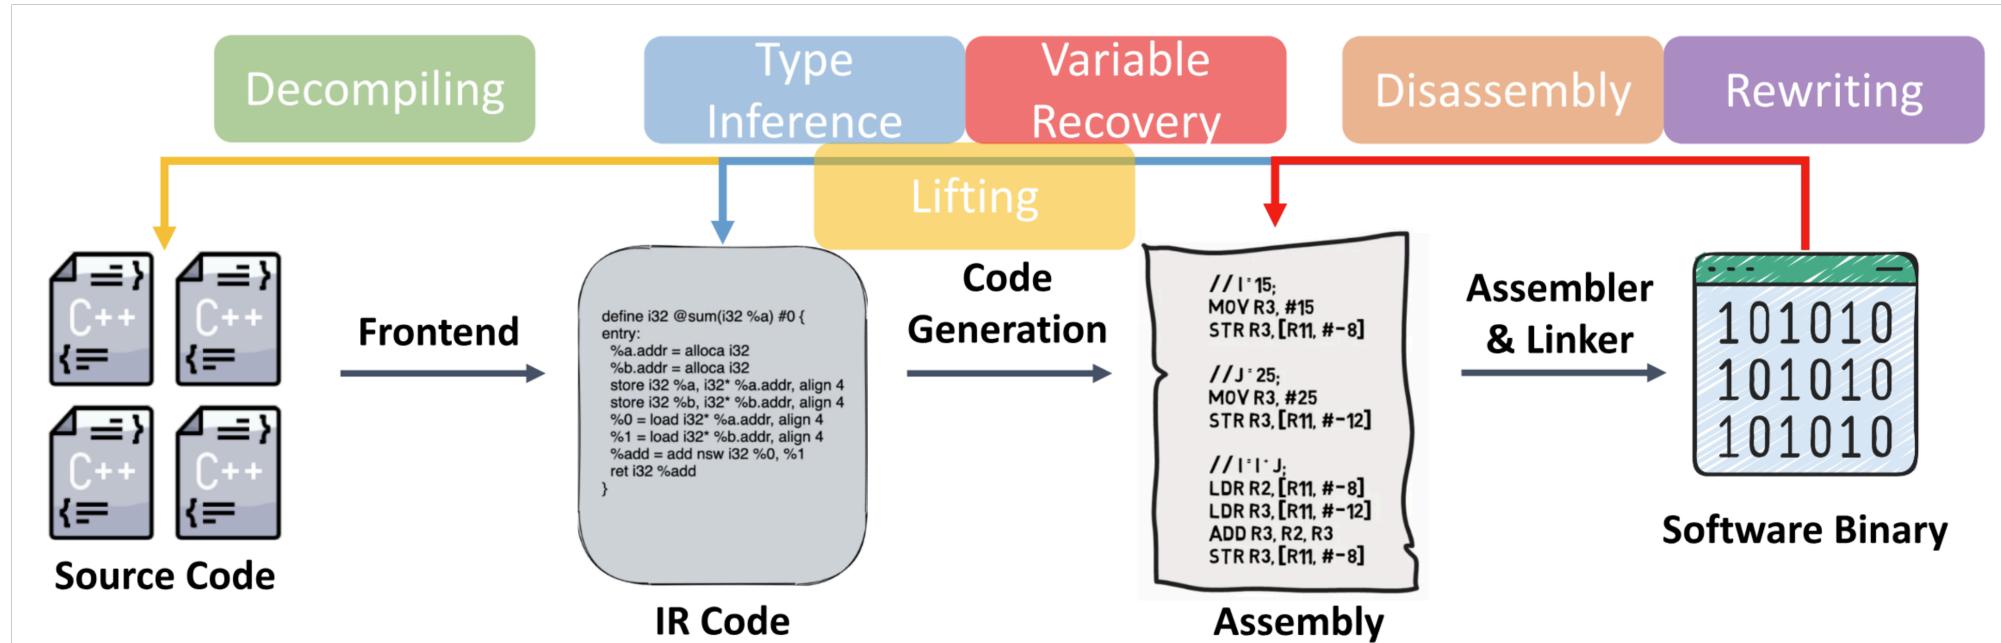


Figure 1: → Compilation ← Reverse Engineering

¹Commercial-Off-The-Shelf, the source code or build environment is not available

Binary Code Analysis

I. Background

- Analyzable binary code formats: Assembly → IR → Source code

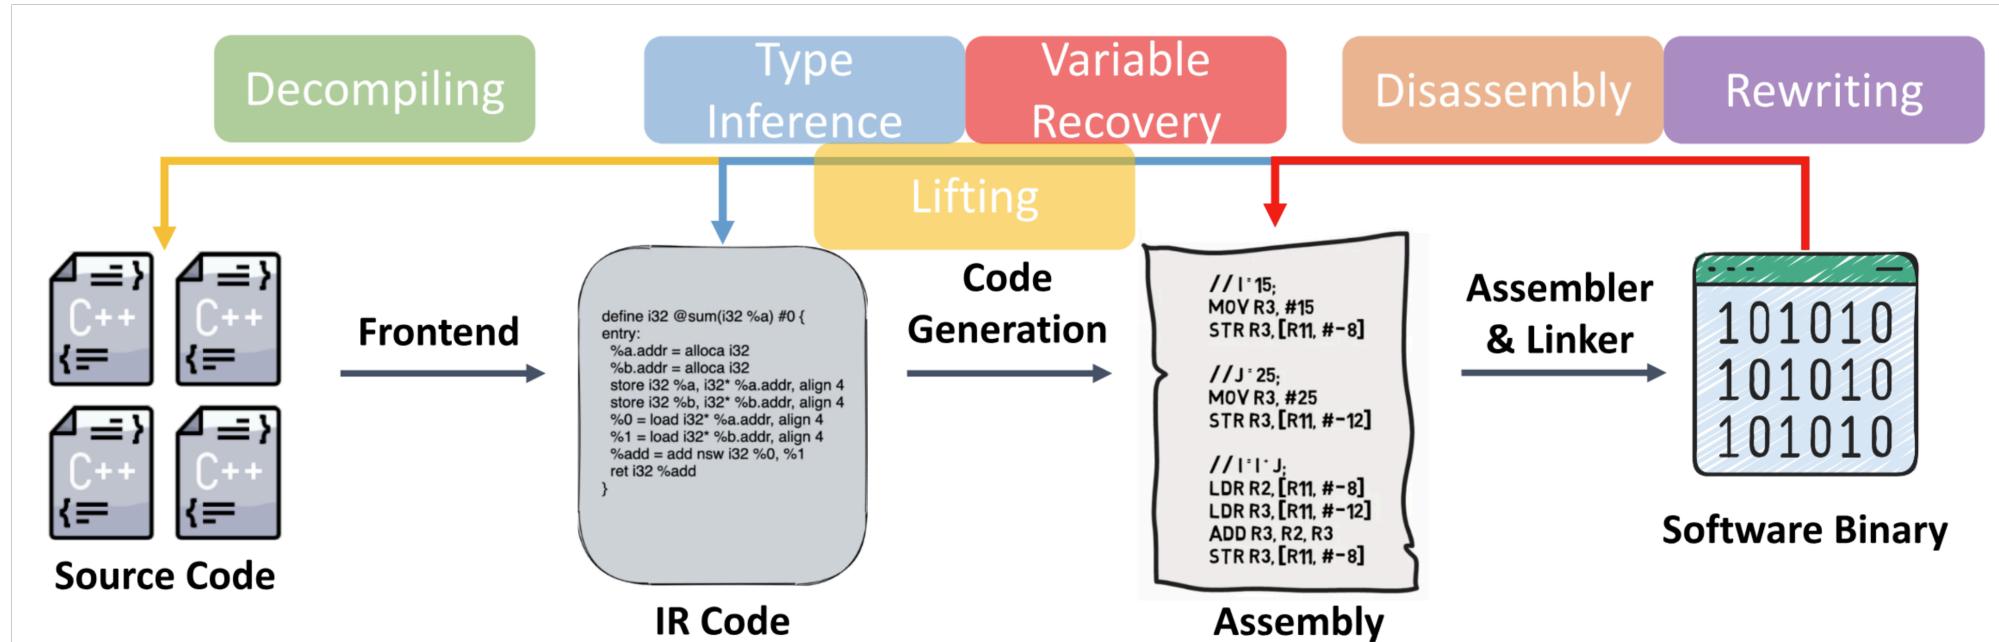


Figure I: → Compilation ← Reverse Engineering

- Binary code analysis is crucial for detecting vulnerabilities in COTS¹ software.

¹Commercial-Off-The-Shelf, the source code or build environment is not available

Challenges

- Traditional Challenges:
 - **Undecidability:** symbol and type information stripped / indirect-call, etc.
 - **Expertise:** ad-hoc human integration and modeling in specific analysis tasks

¹ReSym: Harnessing LLMs to Recover Variable and Data Structure Symbols from Stripped Binaries, CCS '24

- Traditional Challenges:
 - **Undecidability:** symbol and type information stripped / indirect-call, etc.
 - **Expertise:** ad-hoc human integration and modeling in specific analysis tasks
- New challenges in the era of LLM:
 - **Fragmentation:** specialized LLMs are designed for sub-tasks (e.g., variable recovery¹)
 - Hard to merge their knowledge for a comprehensive binary understanding
 - Rely on distinct prompt engineering for effective usage

¹ReSym: Harnessing LLMs to Recover Variable and Data Structure Symbols from Stripped Binaries, CCS '24

- Traditional Challenges:
 - **Undecidability:** symbol and type information stripped / indirect-call, etc.
 - **Expertise:** ad-hoc human integration and modeling in specific analysis tasks
- New challenges in the era of LLM:
 - **Fragmentation:** specialized LLMs are designed for sub-tasks (e.g., variable recovery¹)
 - Hard to merge their knowledge for a comprehensive binary understanding
 - Rely on distinct prompt engineering for effective usage
 - **Hallucination:** general-purpose LLMs struggle to understand large-scale binary code

¹ReSym: Harnessing LLMs to Recover Variable and Data Structure Symbols from Stripped Binaries, CCS '24

- Traditional Challenges:
 - **Undecidability**: symbol and type information stripped / indirect-call, etc.
 - **Expertise**: ad-hoc human integration and modeling in specific analysis tasks
- New challenges in the era of LLM:
 - **Fragmentation**: specialized LLMs are designed for sub-tasks (e.g., variable recovery¹)
 - Hard to merge their knowledge for a comprehensive binary understanding
 - Rely on distinct prompt engineering for effective usage
 - **Hallucination**: general-purpose LLMs struggle to understand large-scale binary code
-  Our direction: let LLM interact with mature binary analysis tools in a guided manner

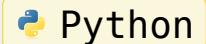
¹ReSym: Harnessing LLMs to Recover Variable and Data Structure Symbols from Stripped Binaries, CCS '24

- Agent¹: anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.

¹Artificial Intelligence: A Modern Approach, 1995

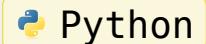
- Agent¹: anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.
- LLM Agent workflow:

```
1 env = Environment()
2 tools = Tools(env)
3 system_prompt = "Goals, constraints, and how to act"
4 while True:
5     action = llm.run(system_prompt + env.state) # agentic 1
6     env.state = tools.run(action)                 # agentic 2
```



- Agent¹: anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.
- LLM Agent workflow:

```
1 env = Environment()
2 tools = Tools(env)
3 system_prompt = "Goals, constraints, and how to act"
4 while True:
5     action = llm.run(system_prompt + env.state) # agentic 1
6     env.state = tools.run(action)                # agentic 2
```



- Agentic:
 1. more informative tool usage
 2. more complex environment, less supervision

¹Artificial Intelligence: A Modern Approach, 1995

- Agent¹: anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.
- LLM Agent workflow:

```
1 env = Environment()
2 tools = Tools(env)
3 system_prompt = "Goals, constraints, and how to act"
4 while True:
5     action = llm.run(system_prompt + env.state) # agentic 1
6     env.state = tools.run(action) # agentic 2
```

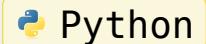
 Python

-  Agentic:
 1. more informative tool usage → help LLM Agent understand binary code
 2. more complex environment, less supervision

¹Artificial Intelligence: A Modern Approach, 1995

- Agent¹: anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.
- LLM Agent workflow:

```
1 env = Environment()
2 tools = Tools(env)
3 system_prompt = "Goals, constraints, and how to act"
4 while True:
5     action = llm.run(system_prompt + env.state) # agentic 1
6     env.state = tools.run(action) # agentic 2
```



- Agentic:
 1. more informative tool usage → help LLM Agent understand binary code
 2. more complex environment, less supervision → guide LLM Agent to explore binary²

¹Artificial Intelligence: A Modern Approach, 1995

²Similar to repo-level exploration (RepoAudit, ICML '25) in source code

Current binary analysis tools for agents¹ have two limitations:

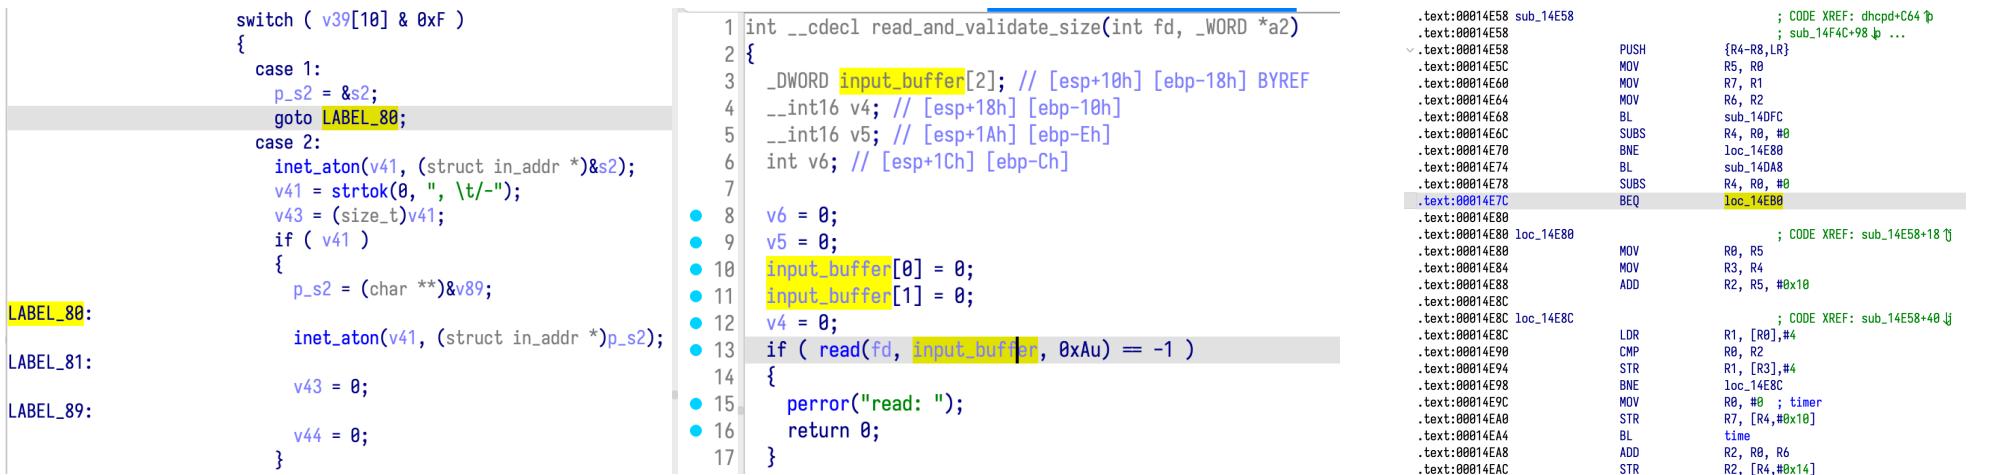
¹EnIGMA: Interactive Tools Substantially Assist LM Agents in Finding Security Vulnerabilities, ICML '25

²We assume perfect disassembly here

Motivation

Current binary analysis tools for agents¹ have two limitations:

I. Co-operation gap between different binary code formats



The screenshot shows a comparison between decompiled code and assembly code for a specific function. The left column, labeled '(a) Decompiled code', contains C-like pseudocode with labels like LABEL_80, LABEL_81, and LABEL_89. The right column, labeled '(b) Assembly code', contains the corresponding assembly language with addresses like .text:00014E58 and various registers like R8, R9, R10, etc. The assembly code is annotated with comments explaining its purpose, such as 'CODE XREF' and memory addresses.

```

(a) Decompiled code
concise (e.g., 50-100 LoC)
fidelity issues: e.g., goto-based control flow, wrong variable length

(b) Assembly code
verbose (e.g., 150-200 LoC)
zero-semantic gap2, less structural

```

Figure 3: Relying on decompiled code and/or assembly code is insufficient

¹EnIGMA: Interactive Tools Substantially Assist LM Agents in Finding Security Vulnerabilities, ICML '25

²We assume perfect disassembly here

Current binary analysis tools for agents have two limitations:

'SV-TrustEval-C: Evaluating Structure and Semantic Reasoning in Large Language Models for Source Code Vulnerability Analysis, IEEE S&P '25

Current binary analysis tools for agents have two limitations:

2. Explore at the function level, with limited guide on intra-procedural control and data flow'

```
{  
    "action": "decompile --function_name FUN_00400df0 rap\\n",  
    "observation": "Decompilation Found!\\n\\nundefined4 FUN_00400df0(void)\\n\\n{\\n    int iVar1;\\n    basic_ostream *pbVar2;\\n    basic_string local_30 [36];\\n    undefined4 local_c;\\n    local_c = 0;\\n    FUN_00400cc0();\\n    std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char>>::basic_string();\\n    ..... /* try { // try from 00400e0d to 00400e7a has its  
    CatchHandler @ 00400e7f */\\n    pbVar2 = std::operator<<((basic_ostream *)std::cout, \"help me\");\\n    std::basic_ostream<char, std::char_traits<char>>::operator<<(\\n        ..... ((basic_ostream<char,  
    std::char_traits<char>> *)pbVar2,\\n        ..... std::endl<char, std::char_traits<char>>);\\n    std::operator>>((basic_istream *)std::cin, local_30);\\n    iVar1 = FUN_00400ce0(local_30);\\n    if (iVar1  
!= 0) {\\n        pbVar2 = std::operator<<((basic_ostream *)std::cout, \"you found me!\");\\n        std::basic_ostream<char, std::char_traits<char>>::operator<<(\\n            ..... ((basic_ostream<char,  
    std::char_traits<char>> *)pbVar2,\\n            ..... std::endl<char, std::char_traits<char>>);\\n    }\\n    local_c = 0;\\n    std::__cxx11::basic_string<char, std::char_traits<char>, std::allocator<char>>::~basic_string\\n    ..... ((basic_string<char, std::char_traits<char>, std::allocator<char>>)local_30);\\n    return local_c;\\n}\\n\\n",  
    "response": "DISCUSSION\\nThe entry function doesn't provide much information. Let's try to decompile the function that's being called as the main function, which is FUN_00400df0.  
\\n\\n```\\ndecompile --function_name FUN_00400df0 rap\\n```",  
    "state": "{\"open_file\": \"/n/a\", \"working_dir\": \"/_home_udiboy_projects_ctf_datasets_LLM_CTF_Database_2020_CSAW-Finals_rev_rap\", \"interactive_session\": \"/n/a\"},  
    \"thought\": \"DISCUSSION\\nThe entry function doesn't provide much information. Let's try to decompile the function that's being called as the main function, which is FUN_00400df0.\\n\\n\",  
    \"execution_time\": 0.10650779701245483  
},
```

Figure 4: Trajectories from EnIGMA, ICML ‘25

★ Cascading effect: fail to understand the tools' output → hallucinate on next action → stuck in exploration → finally might exceed context length

¹SV-TrustEval-C: Evaluating Structure and Semantic Reasoning in Large Language Models for Source Code Vulnerability Analysis, IEEE S&P '25

2. ClearAgent

Overview

A novel Binary Interface for software engineering agent¹ with two tool sets:

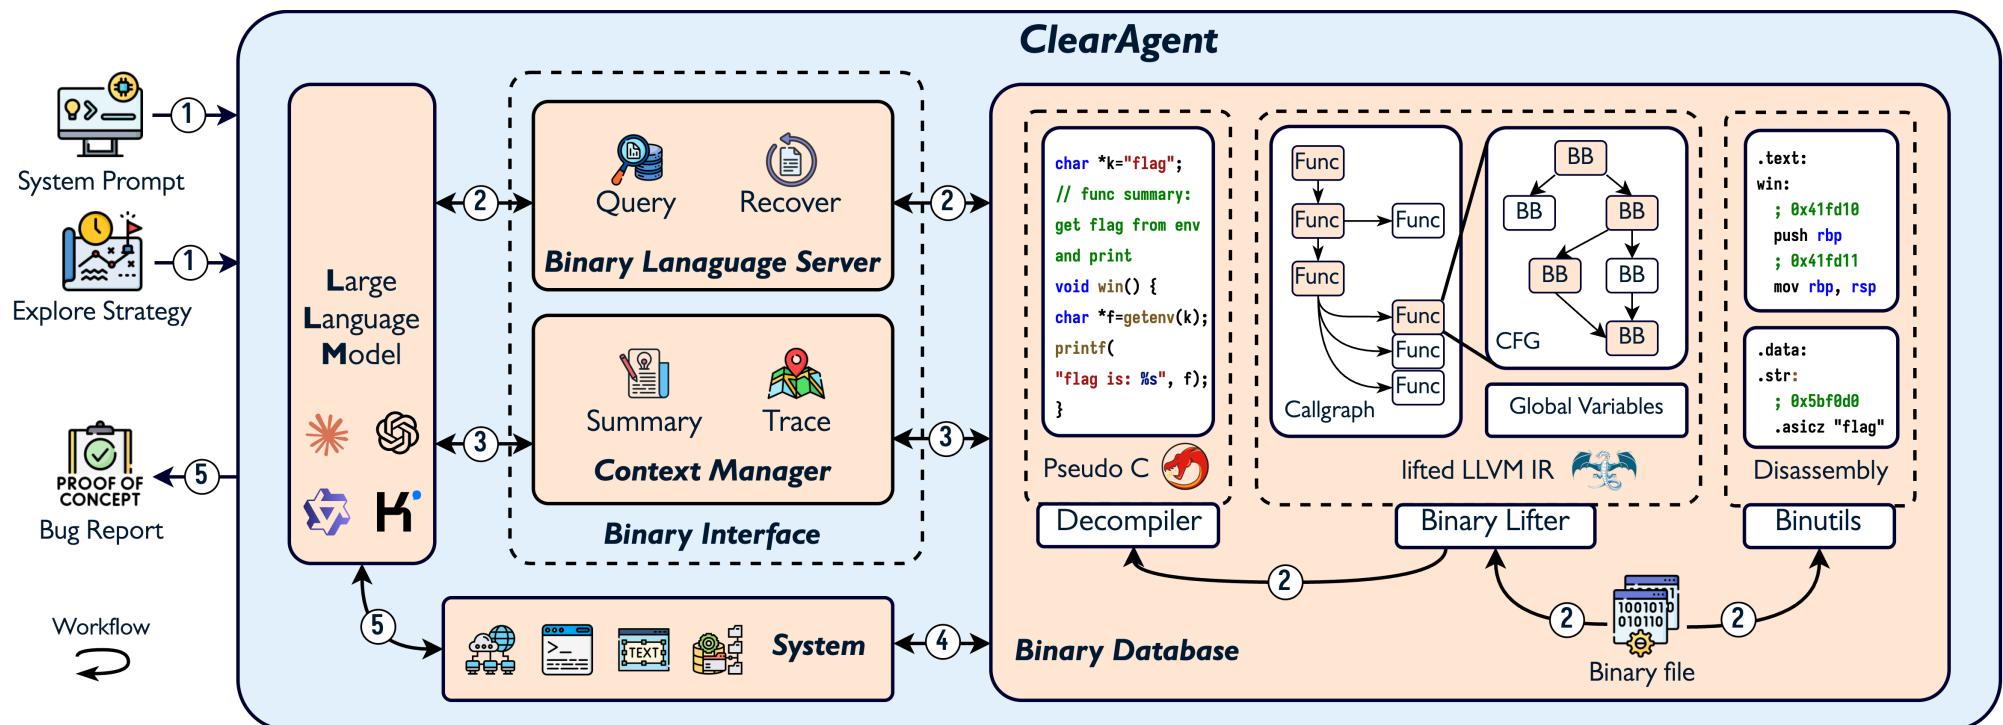
1. **Binary Language Server**: Information retrieval from / update to binary code
2. **Context Manager**: Multi-level exploring binary with summary and feedback

¹Under the Clearblue research project: <https://clearblueinnovations.org/>

Overview

A novel Binary Interface for software engineering agent¹ with two tool sets:

1. **Binary Language Server**: Information retrieval from / update to binary code
2. **Context Manager**: Multi-level exploring binary with summary and feedback



¹Under the Clearblue research project: <https://clearblueinnovations.org/>

- Aims at providing a friendly interface for LLM to understand binary code

¹Plankton: Reconciling Binary Code and Debug Information, ASPLOS '24

²Pinpoint: fast and precise sparse value flow analysis for million lines of code, PLDI '18

- Aims at providing a friendly interface for LLM to understand binary code
- Two-fold friendliness:
 - LLM-friendly: mapping between three formats (Assembly, IR, Pseudo code)
 - Analyzer-friendly: IR¹ as the primary query target, supporting symbolic analysis²

¹Plankton: Reconciling Binary Code and Debug Information, ASPLOS '24

²Pinpoint: fast and precise sparse value flow analysis for million lines of code, PLDI '18

- Aims at providing a friendly interface for LLM to understand binary code
- Two-fold friendliness:
 - LLM-friendly: mapping between three formats (Assembly, IR, Pseudo code)
 - Analyzer-friendly: IR¹ as the primary query target, supporting symbolic analysis²
- Scoped design, narrows down from binary to instruction

Format \ Scope	Binary	Function	Basicblock	Instruction	Global Var
Assembly	✓, target triple	✓, address range	✓, address range	✓, address, size	✓, raw bytes
LLVM IR	✓, symbol	✓, Caller/Callee Signature	✓, Predecessors Successors	✓, Def-Use	✓, User
Pseudo code	✓	✓	✓		

Table I: Queries supported by Binary Language Server, ✓ represents directly querying the code body

¹Plankton: Reconciling Binary Code and Debug Information, ASPLOS '24

²Pinpoint: fast and precise sparse value flow analysis for million lines of code, PLDI '18

Binary Language Server example

2. ClearAgent

How does ClearAgent build the symbolic summary of a function?

```
int32_t do_sta_enrollee_wifi(int32_t arg1)

int32_t $v0 = * __stack_chk_guard
void var_11c
memset(&var_11c, 0, 0x80)
void var_9c
memset(&var_9c, 0, 0x40)
char var_5c
memset(&var_5c, 0, 0x40)
int32_t $v0_1 = getenv(0x461da8) {"wps_sta_enrollee_pin"}
sprintf(0x495850, 0x4628d8, arg1) {"qcawifi.%s.wps_enable"}
int32_t $v0_3 = atoi(uci_safe_get(0x495850))
sprintf(0x495850, 0x4628f0, arg1) {"qcawifi.%s.wps_vap"}
strcpy(&var_9c, uci_safe_get(0x495850))
sprintf(0x495850, 0x462904, &var_9c) {"qcawifi.%s.up"}
int32_t $v0_6 = atoi(uci_safe_get_state(0x495850))
sprintf(0x495850, 0x462914, &var_9c) {"qcawifi.%s.ifname"}
strcpy(&var_5c, uci_safe_get_state(0x495850))

if (access(0x461dd0, 0) == 0) {"/var/tmp/wps_result"}
unlink(0x461dd0) {"/var/tmp/wps_result"}

if ($v0_1 != 0 && $v0_3 != 0 && $v0_6 != 0 && sx.d(var_5c) != 0)
sprintf(&var_11c, 0x462944, &var_5c, $v0_1, __gp) {
    hostapd_cli -i %s wps_pin any %s"
    system(&var_11c)
```

(a) Decompiled code

```
%tmp38 = load [25 x i8]** @global.global_var_49188c, !asm !6193
%tmp39 = getelementptr [25 x i8]* %tmp38, i32 303, i32 17, !asm !6194
%tmp40 = call i8* @getenv(i8* %tmp39), !asm !6194
%tmp41 = ptrtoint i8* %tmp40 to i32, !asm !6194

...
block_416508: ..... ; preds = %block_4164bc
%tmp94 = load [25 x i8]** @global.global_var_49188c, !asm !6233
%tmp95 = getelementptr [25 x i8]* %tmp94, i32 422, i32 14, !asm !6234
%tmp96 = call i32 (i8*, i8*, ...)* @sprintf(i8* %stack.alias.prop18_-284, i8*
%tmp95, i32 %tmp36, i32 %tmp41), !asm !6234
%tmp97 = call i32 @system(i8* %stack.alias.prop18_-284), !asm !6235
```

(b) IR code

```
004164a0 8fbc0010 lw $gp, 0x10($sp) {var_128}
004164a4 12200008 beqz $s1, 0x4164c8
004164a8 8f828cdc lw $v0, -0x7324($gp) {__stack_chk_guard}

004164ac 12600006 beqz $s3, 0x4164c8
004164b0 00000000 nop

004164b4 12400003 beqz $s2, 0x4164c4
004164b8 83a200dc lb $v0, 0xdc($sp) {var_5c}

004164bc 14400012 bnez $v0, 0x416508
004164c0 8f85801c lw $a1, -0x7fe4($gp) {data_49188c}
```

```
00416508 27a4001c addiu $a0, $sp, 0x1c {var_11c}
0041650c 8f998b78 lw $t9, -0x7490($gp) {sprintf}
00416510 27a600dc addiu $a2, $sp, 0xdc {var_5c}
00416514 02203821 move $a3, $s1
00416518 0320f809 jalr $t9
0041651c 24a52944 addiu $a1, $a1, 0x2944 {0x462944, "hostapd_cli"}
00416520 8fbc0010 lw $gp, 0x10($sp) {var_128}
00416524 8f998b6c lw $t9, -0x7494($gp) {system}
00416528 0320f809 jalr $t9
0041652c 27a4001c addiu $a0, $sp, 0x1c {var_11c}
00416530 1000ffe4 b 0x4164c4
00416534 8fbc0010 lw $gp, 0x10($sp) {var_128} {_gp}
```

(c) Assembly code

Figure 7:Three binary code formats of the function do_sta_enrollee_wifi in CVE-2022-46593.

ClearAgent (1) understands this function using decompiled code, then (2) builds the data dependency between %tmp40 and %stack.alias.prop18_-284 while (3) checking the fidelity issue using assembly code

- Guiding LLM Agent to explore the binary at function- and basicblock-level
- Contexts are updated with each query to the binary language server

¹Standardized Operating Procedures (SOPs) in MetaGPT, ICLR '24

²RE-Mind: a First Look Inside the Mind of a Reverse Engineer, Usenix Security '22

Context Manager

- Guiding LLM Agent to explore the binary at function- and basicblock-level
- Contexts are updated with each query to the binary language server

Query \ Scope	Binary	Function	Basicblock
Summary	natural language	natural language, symbolic	natural language, symbolic
Count		visit count of itself, caller/callee	visit count of itself, preds/succs
Trace	visited sub-callgraph	visited sub-CFG	

Table 2: Queries supported by Context Manager

¹Standardized Operating Procedures (SOPs) in MetaGPT, ICLR '24

²RE-Mind: a First Look Inside the Mind of a Reverse Engineer, Usenix Security '22

Context Manager

- Guiding LLM Agent to explore the binary at function- and basicblock-level
- Contexts are updated with each query to the binary language server

Query \ Scope	Binary	Function	Basicblock
Summary	natural language	natural language, symbolic	natural language, symbolic
Count		visit count of itself, caller/callee	visit count of itself, preds/succs
Trace	visited sub-callgraph	visited sub-CFG	

Table 2: Queries supported by Context Manager

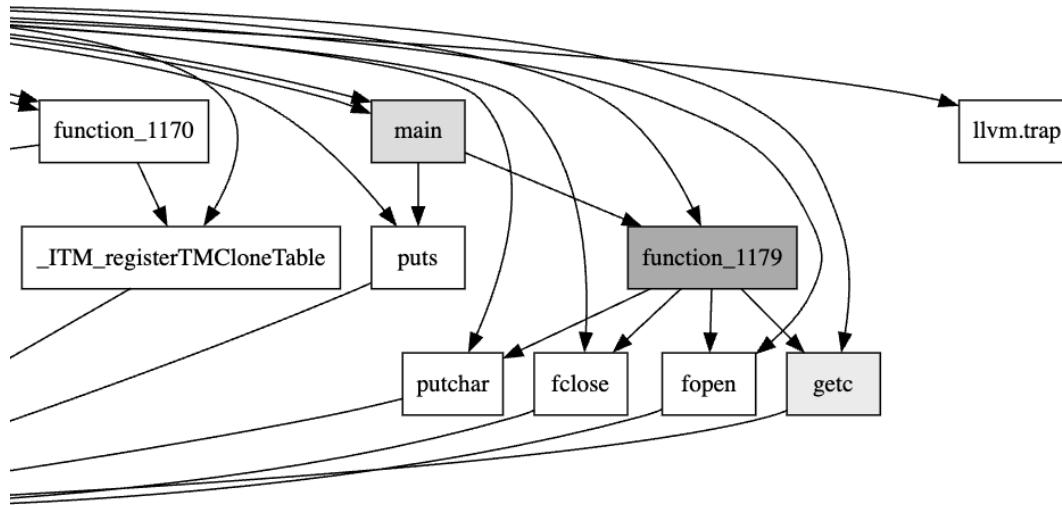
- Collaborating with the exploration strategy
 - Context is managed through phased exploration¹
 - e.g., recovery → analysis → validation
 - Each phase is guided by a pre-defined searching strategy²
 - e.g., Forward / Backward, BFS / DFS

¹Standardized Operating Procedures (SOPs) in MetaGPT, ICLR '24

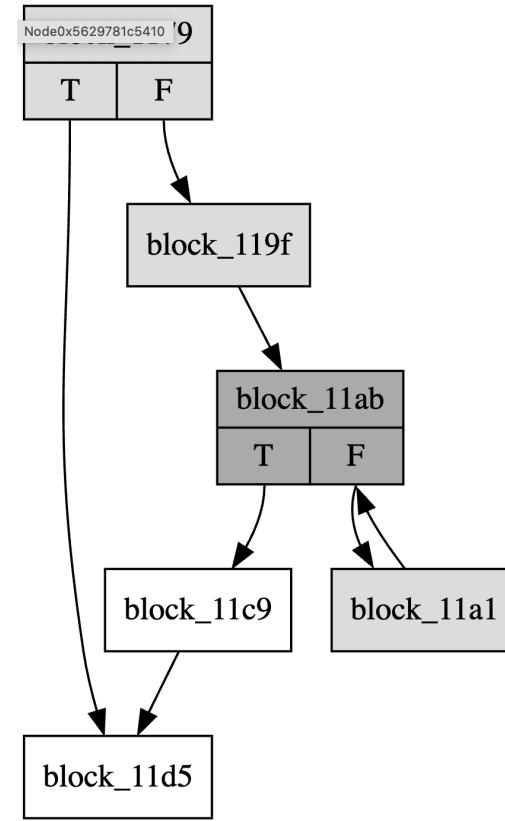
²RE-Mind: a First Look Inside the Mind of a Reverse Engineer, Usenix Security '22

Context Manager example

2. ClearAgent



(a) binary's Callgraph



CFG for 'function_1179' function

(b) function's CFG

Figure 9: Visited node and count formed sub-graph on callgraph and CFG
The darker the color, the more visit

3. Preliminary Evaluation

- Dataset: 90 binary challenges from an entry-level CTF contest¹
 - Only 19/90 can be solved by the SOTA CTF Agents

¹NYU CTF Bench: A Scalable Open-Source Benchmark Dataset for Evaluating LLMs in Offensive Security, NIPS '24

²CRAKEN: Cybersecurity LLM Agent with Knowledge-Based Execution, Arxiv '25

³EnIGMA: Interactive Tools Substantially Assist LM Agents in Finding Security Vulnerabilities, ICML '25

- Dataset: 90 binary challenges from an entry-level CTF contest¹
 - Only 19/90 can be solved by the SOTA CTF Agents
- Metrics:
 - Pwn: Can ClearAgent identify the buggy point and write a PoC script?
 - Reversing: Can ClearAgent understand the logic and get the flag?

¹NYU CTF Bench: A Scalable Open-Source Benchmark Dataset for Evaluating LLMs in Offensive Security, NIPS '24

²CRAKEN: Cybersecurity LLM Agent with Knowledge-Based Execution, Arxiv '25

³EnIGMA: Interactive Tools Substantially Assist LM Agents in Finding Security Vulnerabilities, ICML '25

- Dataset: 90 binary challenges from an entry-level CTF contest¹
 - Only 19/90 can be solved by the SOTA CTF Agents
- Metrics:
 - Pwn: Can ClearAgent identify the buggy point and write a PoC script?
 - Reversing: Can ClearAgent understand the logic and get the flag?

Agent	LLM	Reversing	Pwn	All
ClearAgent	Qwen 3 Plus	7/11	8/8	15/19
CRAKEN ²	Claude 3.5 Sonnet	11/11	6/8	17/19
EnIGMA ³	Claude 3.5 Sonnet	6/11	6/8	12/19

Table 3: Preliminary evaluation on CTF challenges with an early version of ClearAgent (no guidance, only IR)

¹NYU CTF Bench: A Scalable Open-Source Benchmark Dataset for Evaluating LLMs in Offensive Security, NIPS '24

²CRAKEN: Cybersecurity LLM Agent with Knowledge-Based Execution, Arxiv '25

³EnIGMA: Interactive Tools Substantially Assist LM Agents in Finding Security Vulnerabilities, ICML '25

The lifted binary code is inaccurate due to obfuscation or errors in the lifter.

Limitation I - Inaccurate IR

The lifted binary code is inaccurate due to obfuscation or errors in the lifter.

```

loc_400D24:          ; CODE XREF: .text:loc_400D24↑j
    jmp     short near ptr loc_400D24+1
;
dw 48C0h
dq 8B48FFFFFF388563h, 0FFFF30858948F87Dh, 8B48FFFFDF2E8FFh
dq 0C13948FFFFFF308Dh, 858B00000089830Fh, 0F87D8B48FFFFFF38h
dq 89FFFFFF38B56348h, 0FD86E8FFFFFF2C85h, 48B84808BE0FFFFh
dq 0EB04EBC03140EC83h, 0FFFF2C958BB848F5h, 31C231B84806EBFFh
dq 3B84802EBCA31DAh, 4802EBFFFFFF3895h, 0EBFFFFFF388D8BB8h
dq 6348C0FF48B84805h, 943B31C0C303EBC1h, 0C70A74FFFFFF4085h
dq 0FFFFFF3C85h, 0FFF38858B00EB00h, 0FF38858901C083FFh
dq 8BFFFFFF4DE9FFFFh, 0C48148FFFFFF3C85h, 2E66C35D000000E0h
dq 841F0Fh

loc_400D24:          ; CODE XREF: sub_400CE0+F2↓j
nop
inc    eax
movsx  rax, [rbp+i]
mov    rdi, [rbp+string]
mov    [rbp+var_D0], rax
call   __ZNKSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEE6lengthEv

```

Figure 11: Junk bytes in 2020f-rev-rap

Limitation I - Inaccurate IR

The lifted binary code is inaccurate due to obfuscation or errors in the lifter.

```
loc_400D24:                                ; CODE XREF: .text:loc_400D24↑j
    jmp    short near ptr loc_400D24+1
;
dw 48C0h
dq 8B48FFFFFF388563h, 0FFFF30858948F87Dh, 8B48FFFFDF2E8FFh
dq 0C13948FFFFFF308Dh, 858B00000089830Fh, 0F87D8B48FFFFFF38h
dq 89FFFFFF38B56348h, 0FD86E8FFFFFF2C85h, 48B84808BE0FFFFh
dq 0EB04EBC03140EC83h, 0FFFF2C958BB848F5h, 31C231B84806EBFFh
dq 3B84802EBCA31DAh, 4802EBFFFFFF3895h, 0EBFFFFFF388D8BB8h
dq 6348C0FF48B84805h, 943B31C0C303EBC1h, 0C70A74FFFFFF4085h
dq 0FFFFFF3C85h, 0FFFF38858B00EB00h, 0FF38858901C083FFh
dq 8BFFFFFF4DE9FFFFh, 0C48148FFFFFF3C85h, 2E66C35D000000E0h
dq 841F0Fh

loc_400D24:                                ; CODE XREF: sub_400CE0+F2↓j
nop
inc    eax
movsx  rax, [rbp+i]
mov    rdi, [rbp+string]
mov    [rbp+var_D0], rax
call   _ZNKSt7__cxx112basic_stringIcSt11char_traitsIcESaIcEE6lengthEv
```

Figure 11: Junk bytes in 2020f-rev-rap

```
.text:00000000000012A4 ; __ unwind {
.text:00000000000012A4 push   rbp
.text:00000000000012A5 mov    rbp, rsp
.text:00000000000012A8 sub    rsp, 20h
.text:00000000000012AC mov    rax, 6E37625970416742h
.text:00000000000012B6 mov    edx, 44777343h
.text:00000000000012BB mov    qword ptr [rbp+var_20], rax
.text:00000000000012BF mov    qword ptr [rbp+var_20+8], rdx
```

```
define i64 @main(i64 %arg, i8** %arg1) {
block_12a4:
%stack.stack_var_-12 = alloca i64
%stack.stack_var_-16 = alloca i64
%stack.stack_var_-24 = alloca i64
%stack.stack_var_-40 = alloca i64
store i64 7941924604166104898, i64* %stack.stack_var_-40, !asm !0
```

Figure 12: Over-optimized stack variable

Limitation 2 - Insufficient Guidance

3. Preliminary Evaluation

Stuck or hallucinate during inter-procedural exploration, and exceed the round limitation

Limitation 2 - Insufficient Guidance

3. Preliminary Evaluation

Stuck or hallucinate during inter-procedural exploration, and exceed the round limitation

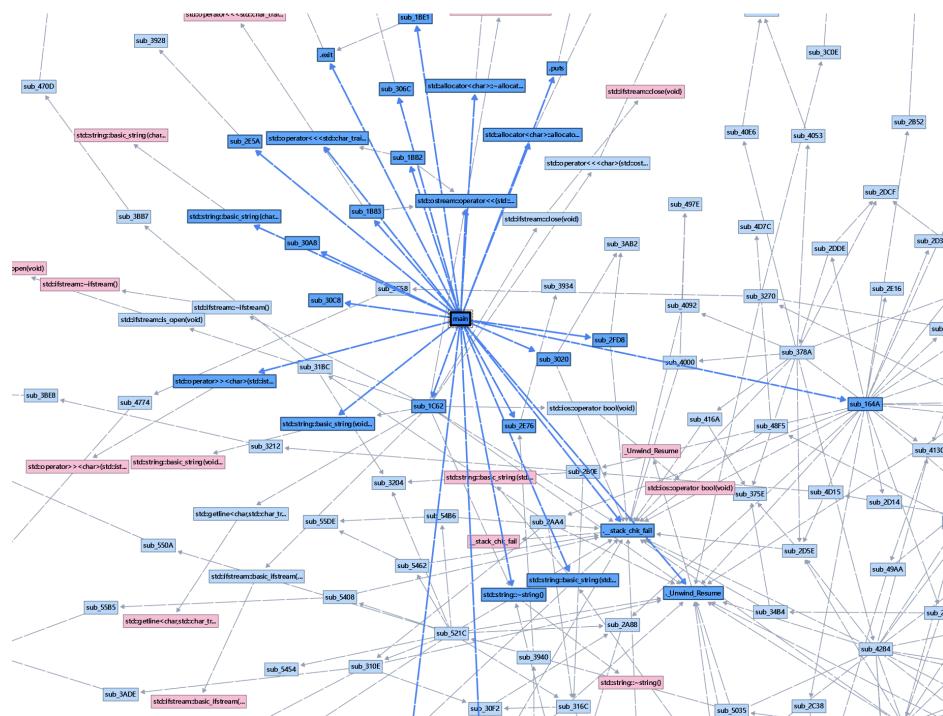


Figure 13: Deep calling context in C++ std functions

Limitation 2 - Insufficient Guidance

Stuck or hallucinate during inter-procedural exploration, and exceed the round limitation

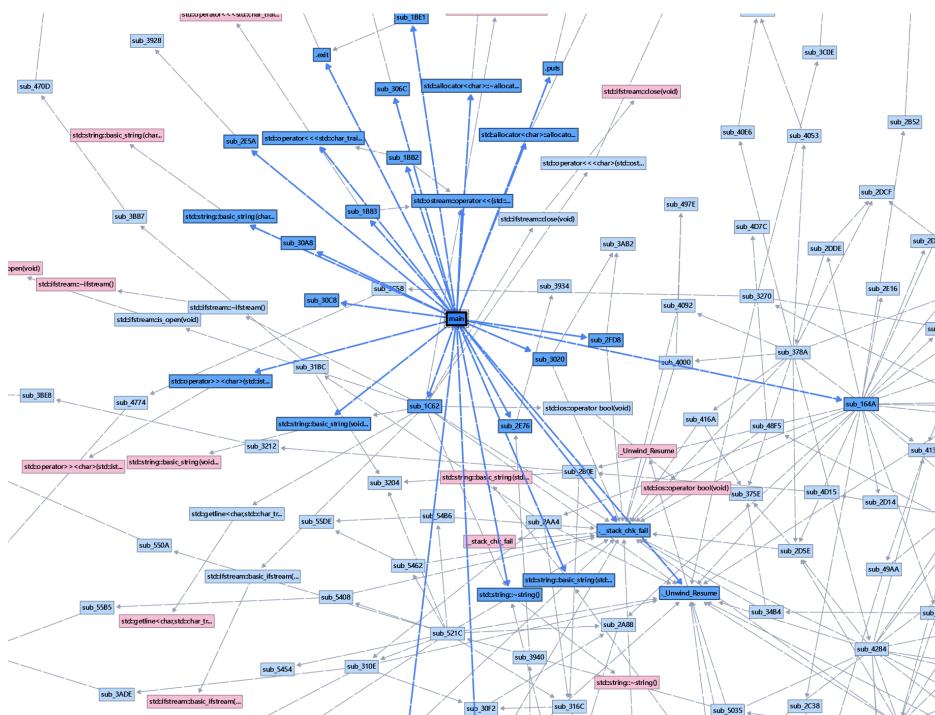


Figure 13: Deep calling context in C++ std functions

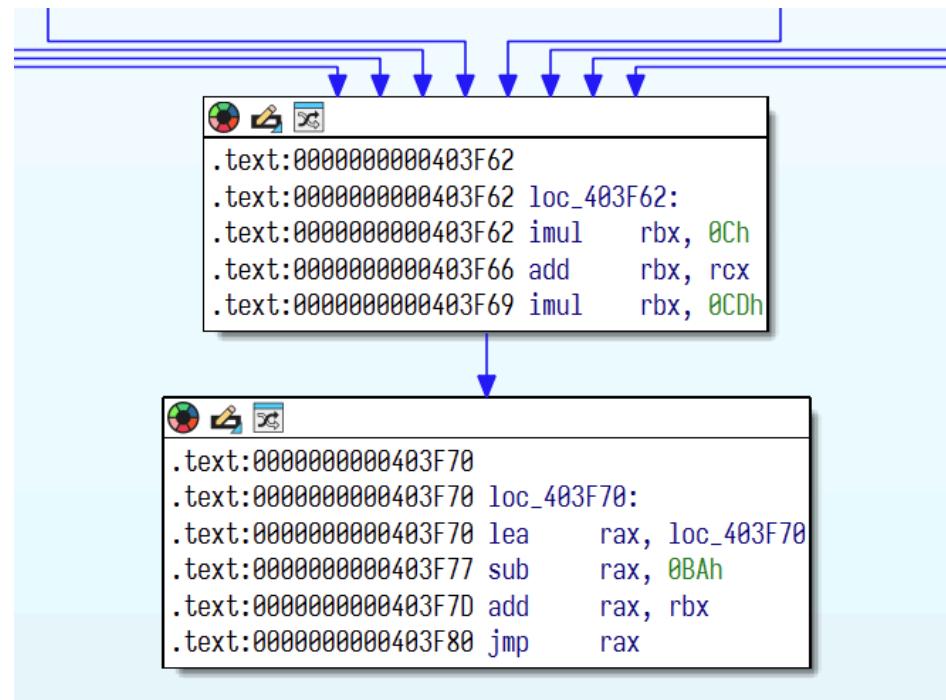
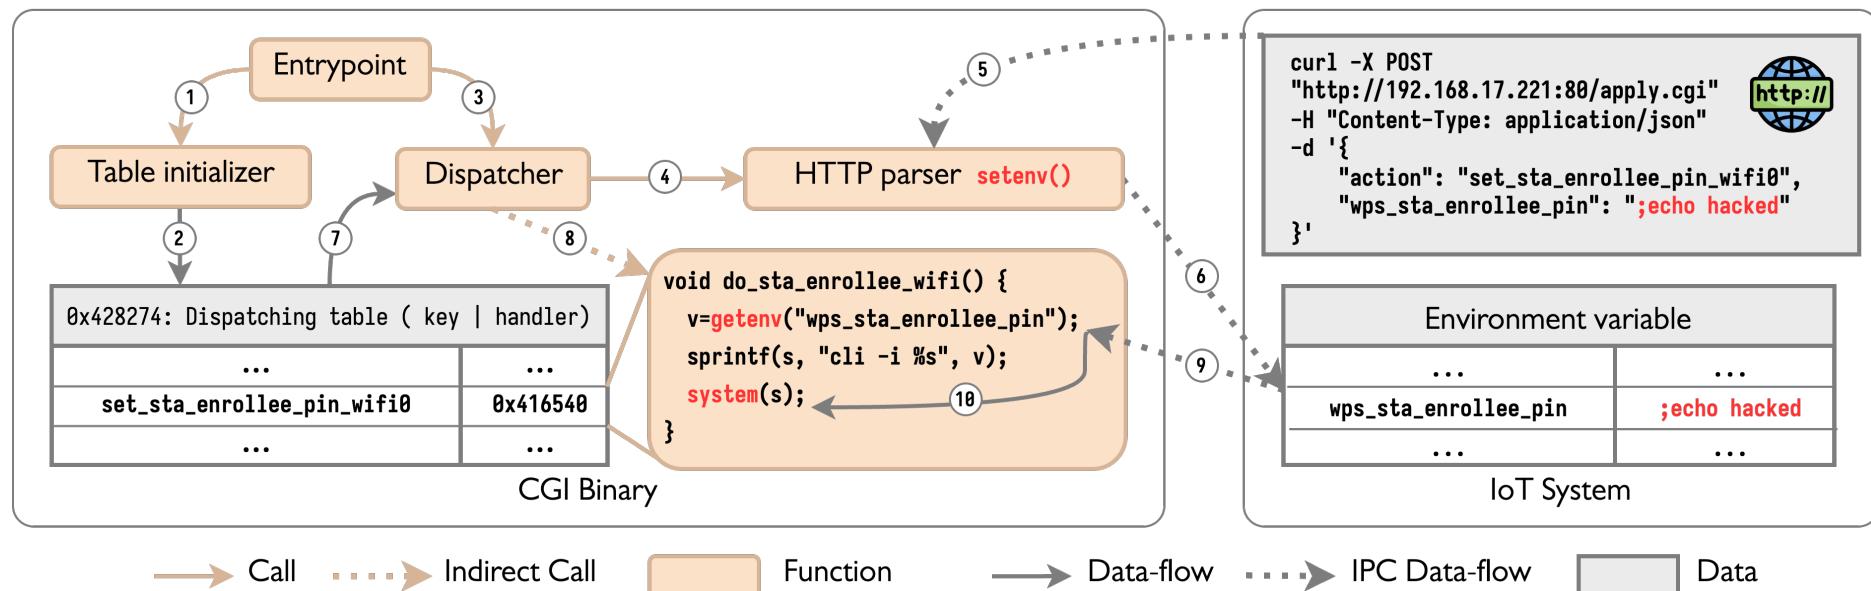


Figure 14: Indirect call targets are missed by static analysis

RQ2 - Context Manager

- Demo: CGI Binary in TRENDnet IoT firmware¹
 - **Recovery:** find and recover the dispatching table using DFS/Forward strategy
 - **Analysis:** analyze each handler using BFS/Forward strategy at the indirect callsite
 - **Validation:** find the setenv operation using DFS/Backward strategy



¹<https://nvd.nist.gov/vuln/detail/CVE-2022-46593>

4. Future Work

Directions

We envision the following three future directions:

- Agentic IR refinement: fixing fidelity issues in lifted IR, such as type recovery¹
- Multi-binary analysis²: cross-binary vulnerability in IoT firmware
- Huge binary analysis (~1GB): pruning uninteresting functions in rigorous static analysis

¹Think of IDA Pro: <https://hex-rays.com/blog/igors-tip-of-the-week-42-renaming-and-retyping-in-the-decompiler>

²Karonte: Detecting Insecure Multi-binary Interactions in Embedded Firmware, IEEE S&P '20

Any Question